

```
In [10]: import pandas as pd  
  
df=pd.read_csv('dataset/fuel_consumption_dataset.csv')  
  
# df.info()  
# df.head()  
# df.tail()  
df.describe()
```

Out[10]:

	MODELYEAR	ENGINESIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY	FUELCONSUMPTION_COMB	FUELCONSUMPTIC
count	1067.0	1067.000000	1067.000000	1067.000000	1067.000000	1067.000000	1067.000000
mean	2014.0	3.346298	5.794752	13.296532	9.474602	11.580881	
std	0.0	1.415895	1.797447	4.101253	2.794510	3.485595	
min	2014.0	1.000000	3.000000	4.600000	4.900000	4.700000	
25%	2014.0	2.000000	4.000000	10.250000	7.500000	9.000000	
50%	2014.0	3.400000	6.000000	12.600000	8.800000	10.900000	
75%	2014.0	4.300000	8.000000	15.550000	10.850000	13.350000	
max	2014.0	8.400000	12.000000	30.200000	20.500000	25.800000	



```
In [8]: import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# Select features (X) and target (y)
X = df[['ENGINESIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']]
y = df['CO2EMISSIONS']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Metrics
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
# rmse = np.sqrt(mse)


print(" Model Performance Metrics")
print("R2 Score:", r2)
print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
# print("Root Mean Squared Error (RMSE):", rmse)

# Plot Actual vs Predicted
plt.figure(figsize=(7, 6))
plt.scatter(y_test, y_pred, alpha=0.6, color='blue', label="Predicted vs Actual")
plt.plot([y.min(), y.max()], [y.min(), y.max()], color='red', linestyle='--', label="Perfect Prediction")
plt.xlabel("Actual CO2 Emissions")
plt.ylabel("Predicted CO2 Emissions")
plt.title("Actual vs Predicted CO2 Emissions")
plt.legend()
```

```
plt.grid(True, linestyle="--", alpha=0.6)
plt.show()
```

#### Model Performance Metrics

R<sup>2</sup> Score (Accuracy): 0.8759705206914069  
Mean Absolute Error (MAE): 16.7215939835165  
Mean Squared Error (MSE): 512.8551370148303





```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

# Load dataset
file_path = "dataset/fuel_consumption_dataset.csv"
df = pd.read_csv(file_path)

# Create binary target: High (1) if CO2 > mean, else Low (0)
threshold = df['CO2EMISSIONS'].mean()
df['HighEmission'] = (df['CO2EMISSIONS'] > threshold).astype(int)

# Features and target
X = df[['ENGINESIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']]
y = df['HighEmission']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Logistic Regression model
model = LogisticRegression(max_iter=2000)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Metrics
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))

# Detailed classification report
print("\n Classification Report:")
print(classification_report(y_test, y_pred))

# Plot Actual vs Predicted
plt.figure(figsize=(7, 6))
```

```
plt.scatter(y_test, y_pred, alpha=0.6, color='blue', label="Predicted vs Actual")
plt.plot([y.min(), y.max()], [y.min(), y.max()], color='red', linestyle='--', label="Perfect Prediction")
plt.xlabel("Actual CO2 Emissions")
plt.ylabel("Predicted CO2 Emissions")
plt.title("Actual vs Predicted CO2 Emissions")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.6)
plt.show()
```

Accuracy: 0.9532710280373832


Precision: 0.9473684210526315

Recall: 0.9473684210526315

F1 Score: 0.9473684210526315

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.96	0.96	119
1	0.95	0.95	0.95	95
accuracy			0.95	214
macro avg	0.95	0.95	0.95	214
weighted avg	0.95	0.95	0.95	214



In [12]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1067 entries, 0 to 1066
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   MODELYEAR        1067 non-null    int64  
 1   MAKE              1067 non-null    object  
 2   MODEL             1067 non-null    object  
 3   VEHICLECLASS     1067 non-null    object  
 4   ENGINESIZE        1067 non-null    float64 
 5   CYLINDERS         1067 non-null    int64  
 6   TRANSMISSION      1067 non-null    object  
 7   FUELTYPE          1067 non-null    object  
 8   FUELCONSUMPTION_CITY 1067 non-null    float64 
 9   FUELCONSUMPTION_HWY   1067 non-null    float64 
 10  FUELCONSUMPTION_COMB 1067 non-null    float64 
 11  FUELCONSUMPTION_COMB MPG 1067 non-null    int64  
 12  CO2EMISSIONS      1067 non-null    int64  
dtypes: float64(4), int64(4), object(5)
memory usage: 108.5+ KB
```

1. Display the first 10 rows of the dataset
2. Show all unique vehicle classes
3. Find the average CO2 emissions grouped by fuel type
4. List all cars from the year 2018
5. Show top 5 makes with the highest average engine size
6. Count the number of cars per cylinder type
7. Filter cars where CO2 emissions are greater than 300
8. Find the vehicle with the maximum highway fuel consumption
9. Group by make and model, and find the average city fuel consumption
10. Find the correlation between engine size and CO2 emissions



```
In [18]: import pandas as pd
```

```
# Load the dataset  
df=pd.read_csv('dataset/fuel_consumption_dataset.csv')
```

```
# 1. Display the first 10 rows of the dataset  
query1 = df.head(10)
```

```
# 2. Show all unique vehicle classes  
query2 = df['VEHICLECLASS'].unique()
```

```
# 3. Find the average CO2 emissions grouped by fuel type  
query3 = df.groupby('FUELTYPE')['CO2EMISSIONS'].mean()
```

```
# 4. List all cars from the year 2018  
query4 = df[df['MODELYEAR'] == 2018]
```

```
# 5. Show top 5 makes with the highest average engine size  
query5 = df.groupby('MAKE')['ENGINESIZE'].mean().sort_values(ascending=False).head(5)
```

```
# 6. Count the number of cars per cylinder type  
query6 = df['CYLINDERS'].value_counts()
```

```
# 7. Filter cars where CO2 emissions are greater than 300  
query7 = df[df['CO2EMISSIONS'] > 300]
```

```
# 8. Find the vehicle with the maximum highway fuel consumption  
query8 = df.loc[df['FUELCONSUMPTION_HWY'].idxmax()]
```

```
# 9. Group by make and model, and find the average city fuel consumption  
query9 = df.groupby(['MAKE', 'MODEL'])['FUELCONSUMPTION_CITY'].mean()
```

```
# 10. Find the correlation between engine size and CO2 emissions  
query10 = df[['ENGINESIZE', 'CO2EMISSIONS']].corr().iloc[0,1]
```

```
# Example to print one of the queries
```

```
print(query5)
```

```
MAKE
SRT          8.400000
ROLLS-ROYCE 6.657143
LAMBORGHINI 5.633333
BENTLEY     5.350000
ASTON MARTIN 5.214286
Name: ENGINESIZE, dtype: float64
```